

ENHANCING RELATIONAL MODELS WITH GRAPH PROCESSING IN SQL SERVER 2017

TERRY MCCANN

@SQLSHARK





*We enable our clients
make sense of their data*



DATA ARCHITECTURE



MANAGED SERVICES



DATA STRATEGY



DATA SCIENCE



DATA MANAGEMENT



DATA ANALYTICS



Shout it out!!

Where to find the
slides and *demos*


bit.ly/2uHQS74
Github/SQLShark

Terry McCann 2017

SQL Server 2017 Graph data processing. An introduction.

21. April 2017 terry MCCANN

(2)



- Supports **Graph objects** & **Graph queries** to analyze complex relationships
- **Adaptive query processing** learns & optimizes for unparalleled performance
- **Python** and **R** support
- Advanced ML + **Deep Learning on GPUs**

Industry-leading performance on the most secure data platform, with built-in intelligence for all your data

On the 19th of April 2017 at the Microsoft Data AMP, Microsoft announced SQL Server 2017 and a few new advanced analytics features (slide above). You can watch the full AMP here <https://www.microsoft.com/en-us/sql-server/data-amp>. One of the announcements related to the new support of Graph objects. There have been rumblings over the past few years of Microsoft working on a Graph engine. Project Trinity appeared to be just that <https://www.microsoft.com/en-us/research/project/trinity/>. Trinity never really made much of an impact, possibly due to other vendors having an easier to use product. Hopefully Microsoft are changing this by introducing graph querying directly in to the engine for SQL Server 2017. In this blog post I will look at the upcoming SQL Graph the graph processing engine in SQL Server 2017.

What is a Graph Database?


Relational databases are fantastic at answering many types of queries. There are developers who can write a


In this session we will look to understand what is ***Graph Processing***?

We will explore the history, the theory, the use cases, ***NODE, EDGE & MATCH***.

We will understand why Graph is a ***big deal***, and why it will ***change*** how we work with highly connected data.

What is a Graph Database?



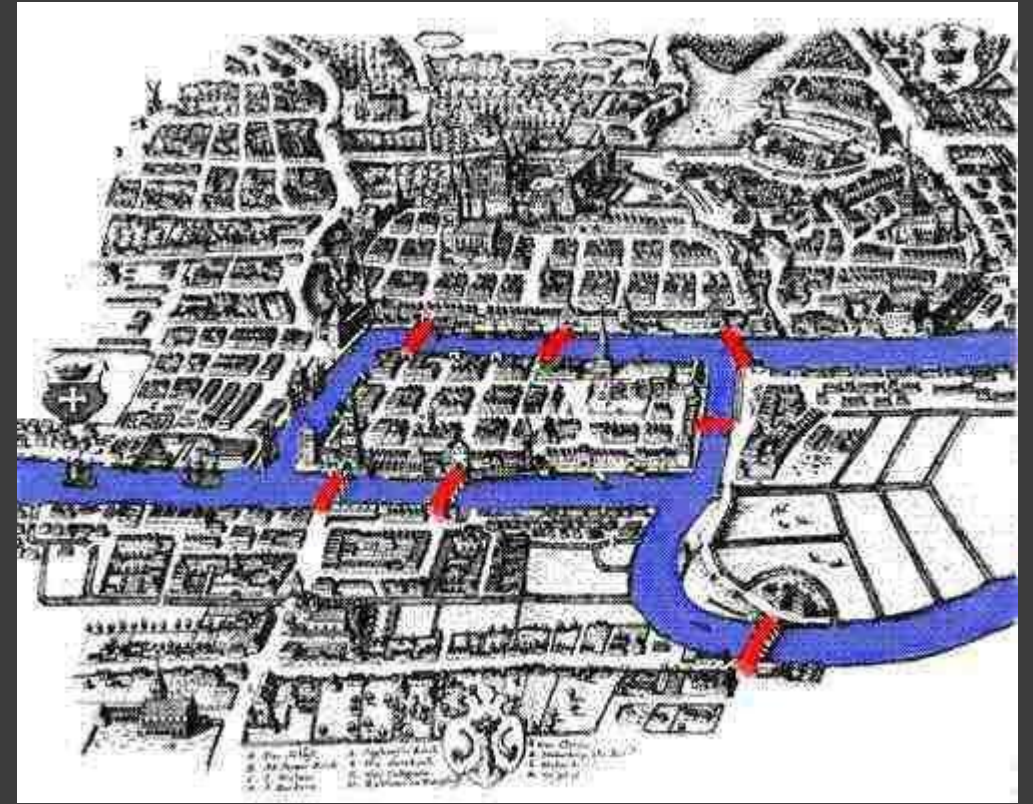


The History of Graph Databases

Königsberg Bridge Problem

Königsberg Bridge Problem









“Which route would allow you to cross all 7 bridges without crossing any more than once”



Take a minute to work it out



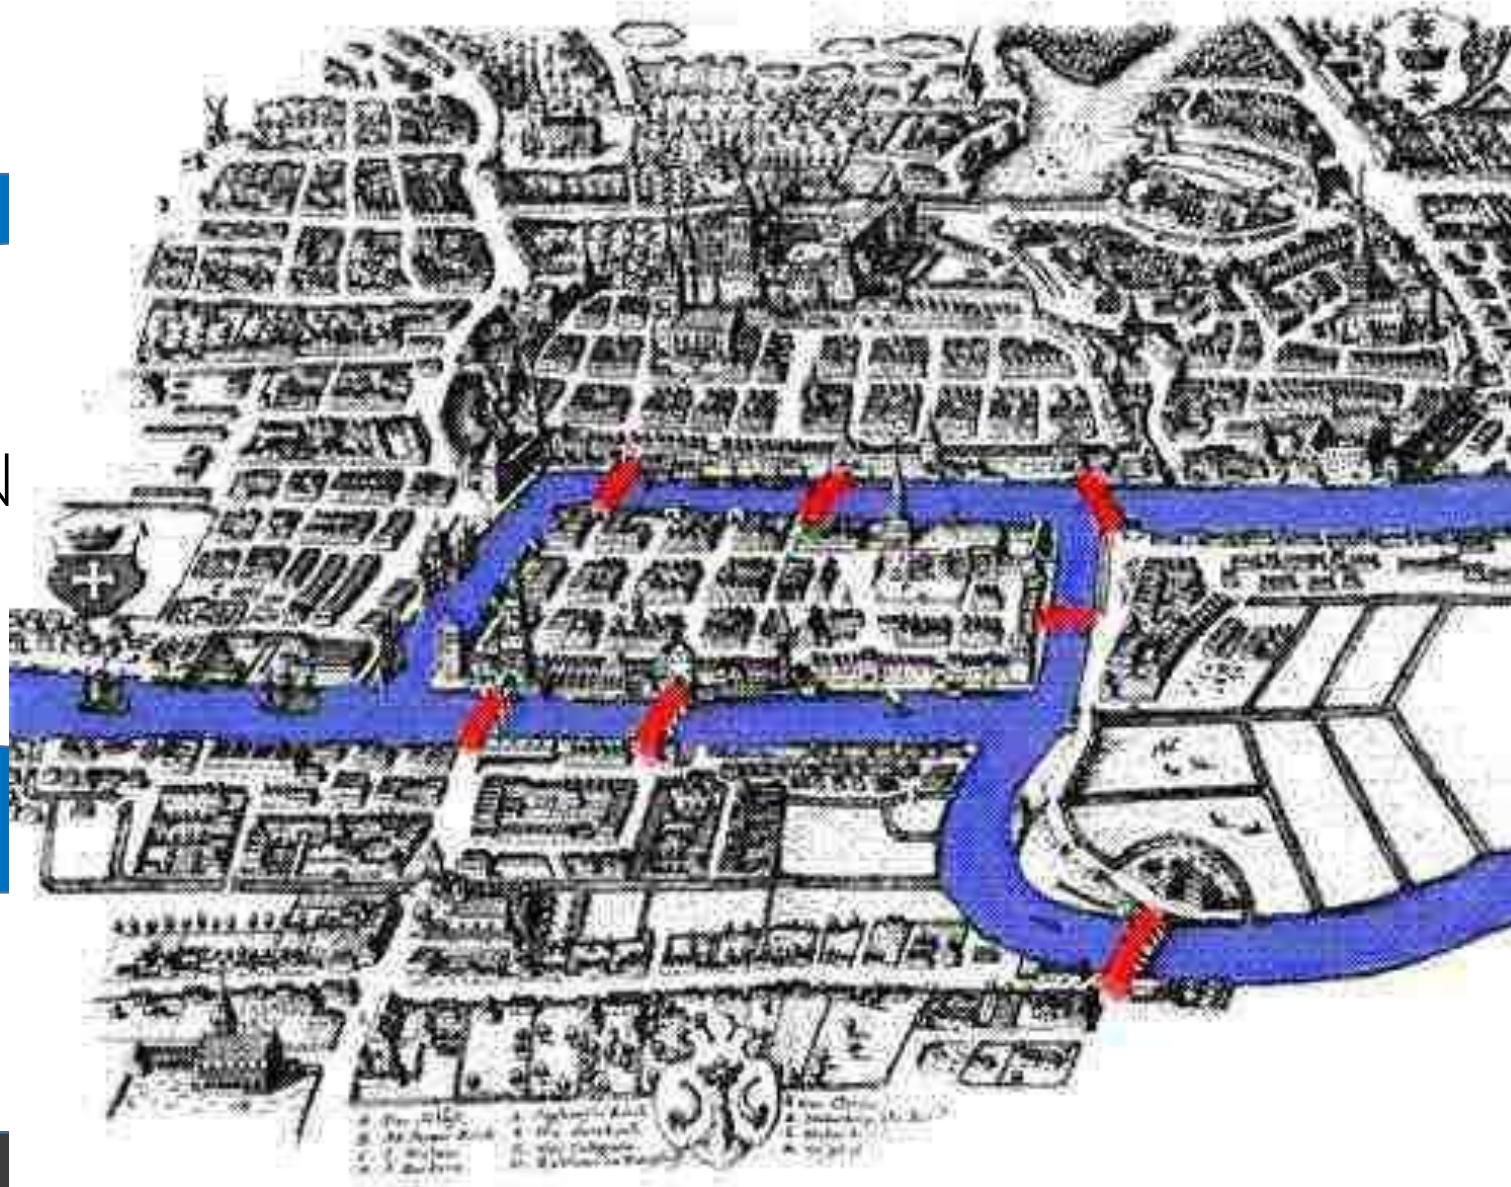
Solved it?




Leonhard Euler



ISLAN



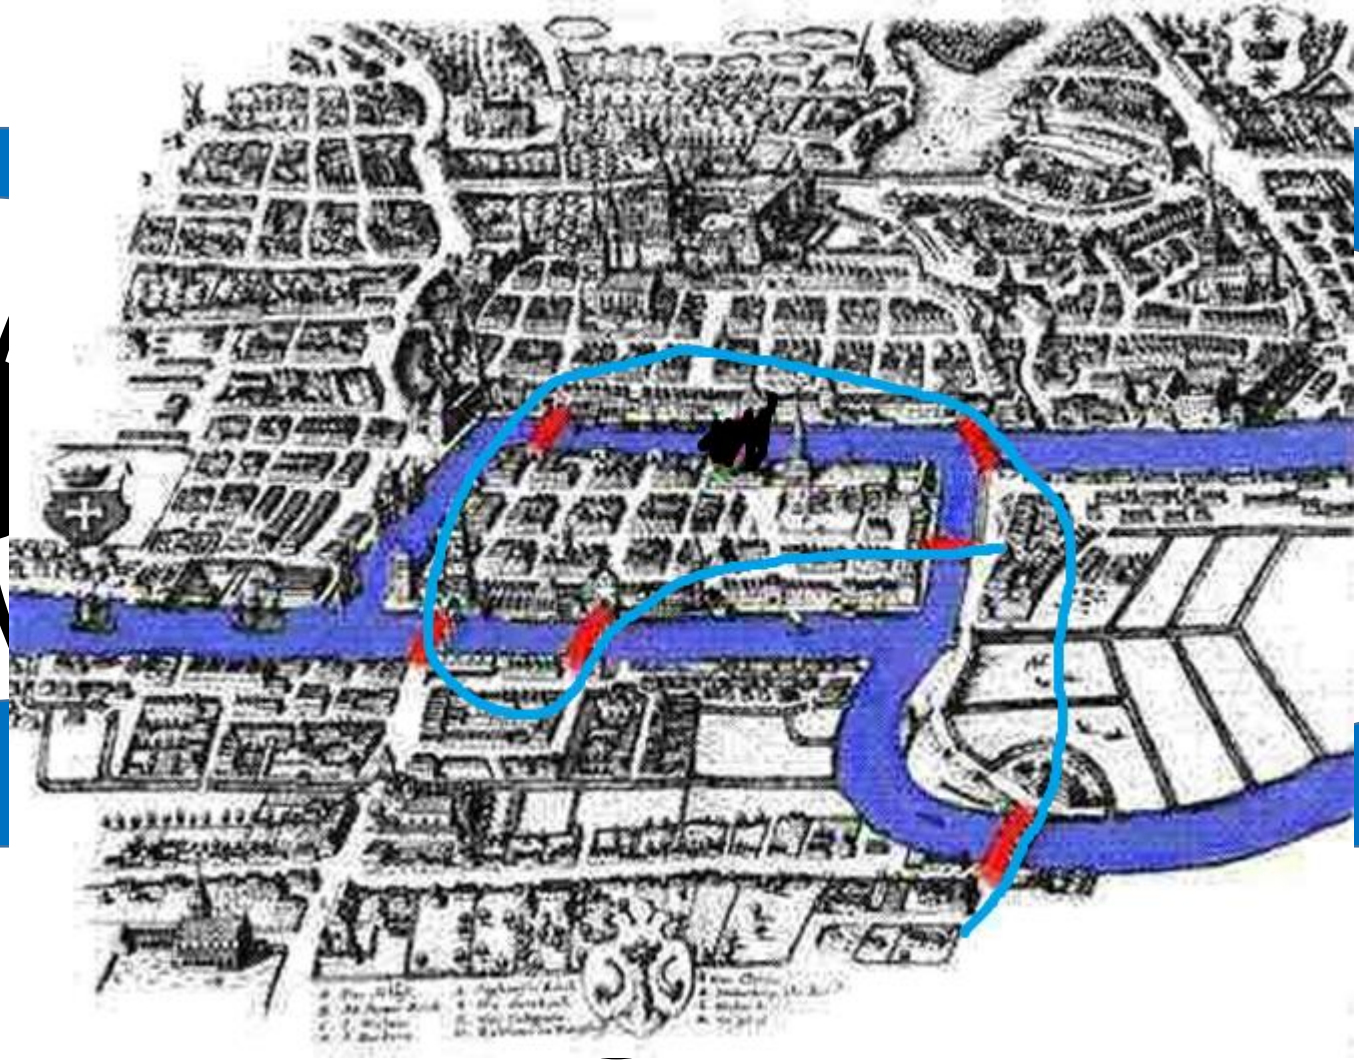


NODE


NODE

4


NODE



Leonhard Euler



Graph Theory



NODE – An entity in a graph

EDGE – Relationship between NODEs

What is a Graph Database?

In computing, a **graph database** is a database that uses graph structures for semantic queries with **nodes**, **edges** and **properties** to represent and store data.

Graph databases, by design, allow **simple** and **fast** retrieval of **complex** hierarchical structures that are **difficult** to model in relational systems


You cannot do that in a Relational Database



The *problem* with relational databases

Why *relational* database?

A Simple Example



Impedance Mismatch

Date Invoice #

CUSTOMER INFORMATION

Name
 Contact
 Info
 E-Mail

Terms of Purchase
☐ PAID #

Salesperson

Product #	Description	Qty	Price	Extended	Tax Exempt
1	Computer	1	850	850.00	<input type="checkbox"/>
2	LCD Monitor	1	250	250.00	<input type="checkbox"/>
3	Mouse	1	30	30.00	<input type="checkbox"/>
4					<input type="checkbox"/>
5					<input type="checkbox"/>
6					<input type="checkbox"/>
7					<input type="checkbox"/>
8					<input type="checkbox"/>
9					<input type="checkbox"/>
10					<input type="checkbox"/>
11					<input type="checkbox"/>
12					<input type="checkbox"/>
13					<input type="checkbox"/>
14					<input type="checkbox"/>
15					<input type="checkbox"/>

Message

SUBTOTAL 1,130.00
Discount 0.00
SUBTOTAL 1,130.00
Shipping 0.00
Sales Tax 56.50
TOTAL 1,186.50

CUSTOMER

<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>

INVOICE

<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>

PRODUCT

<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>

MESSAGES

<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>
<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>

Complex to model and store relationships

Performance *degrades* with more data

Queries can be *long* and *complex*

What is complexity in a database?

Complexity=
 $f(\text{size}, \text{variable structure}, \text{connectedness})$

MapReduce: Si

Bigtable: A Distributed Storage System for Structured Data

Fay Chang, Jeffrey Dean, Sanjay Ghemawat, V
Mike Burrows, Tushar Chandra, And

{fay,jeff,sanjay,wilsonh,kerr,m3h,tushar,

Google, Inc

Abstract

MapReduce is a programming model for processing large data sets. Users specify a *map* function that generates a set of intermediate key/value pairs, and a *reduce* function that merges the values associated with the same intermediate key. In this paper, we describe how real world tasks are expressible in this model.

Programs written in this functional model are automatically parallelized and executed on a large number of commodity machines. The run-time system handles the details of partitioning the input data, scheduling the program's execution across a set of machines, handling machine failures, and managing the required resources.

Abstract

Bigtable is a distributed storage system for managing structured data that is designed to scale to a very large size: petabytes of data across thousands of commodity servers. Many projects at Google store data in Bigtable, including web indexing, Google Earth, and Google Finance. These applications place very different demands on Bigtable, both in terms of data size (from URLs to web pages to satellite imagery) and latency requirements (from backend bulk processing to real-time data serving). Despite these varied demands, Bigtable has successfully provided a flexible, high-performance solution for all of these Google products. In this paper we describe the simple data model provided by Bigtable, which gives clients dynamic control over data layout and format, and we describe the design and implementation of Bigtable.

ach
pro
doc
pro
dyn
low
dat
dex
stri
alth
ture
can
cho
ram
dat
S
Sec

Dynamo: Amazon's Highly Available Key-value Store

Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall and Werner Vogels

Amazon.com

ABSTRACT

Reliability at massive scale is one of the biggest challenges we face at Amazon.com, one of the largest e-commerce operations in the world; even the slightest outage has significant financial consequences and impacts customer trust. The Amazon.com platform, which provides services for many web sites worldwide, is implemented on top of an infrastructure of tens of thousands of servers and network components located in many datacenters around the world. At this scale, small and large components fail continuously and the way persistent state is managed in the face of these failures drives the reliability and scalability of the software systems.

This paper presents the design and implementation of Dynamo, a highly available key-value storage system that some of Amazon's core services use to provide an "always-on" experience. To achieve this level of availability, Dynamo sacrifices consistency under certain failure scenarios. It makes extensive use of object versioning and application-assisted conflict resolution in a manner that provides a novel interface for developers to use.

Categories and Subject Descriptors

D.4.2 [Operating Systems]: Storage Management; D.4.5 [Operating Systems]: Reliability; D.4.2 [Operating Systems]: Performance.

General Terms

Algorithms, Management, Measurement, Performance, Design, Reliability.

1. INTRODUCTION

Amazon runs a world-wide e-commerce platform that serves tens of millions of customers in more than a dozen countries. It has thousands of

One of the lessons our organization has learned from operating Amazon's platform is that the reliability and scalability of a system is dependent on how its application state is managed. Amazon uses a highly decentralized, loosely coupled, service-oriented architecture consisting of hundreds of services. In this environment there is a particular need for storage technologies that are always available. For example, customers should be able to view and add items to their shopping cart even if disks fail, network routes are flapping, or data centers are being destroyed by tornados. Therefore, the service responsible for managing shopping carts requires that it can always write to and read from its data store, and that its data needs to be available across multiple data centers.

Dealing with failures in an infrastructure comprised of millions of components is our standard mode of operation; there are always a small but significant number of server and network components that are failing at any given time. As such Amazon's software systems need to be constructed in a manner that treats failure as the normal case without impacting availability or performance.

To meet the reliability and scaling needs, Amazon has developed a number of storage technologies, of which the Amazon Simple Storage Service (also available outside of Amazon and known as Amazon S3), is probably the best known. This paper presents the design and implementation of Dynamo, another highly available and scalable distributed data store built for Amazon's platform. Dynamo is used to manage the state of services that have high reliability requirements and need tight control over the tradeoffs between availability, consistency, cost-effectiveness, and performance. Amazon's platform has a very diverse set of applications with different storage requirements. A select set of

Polyglot Persistence

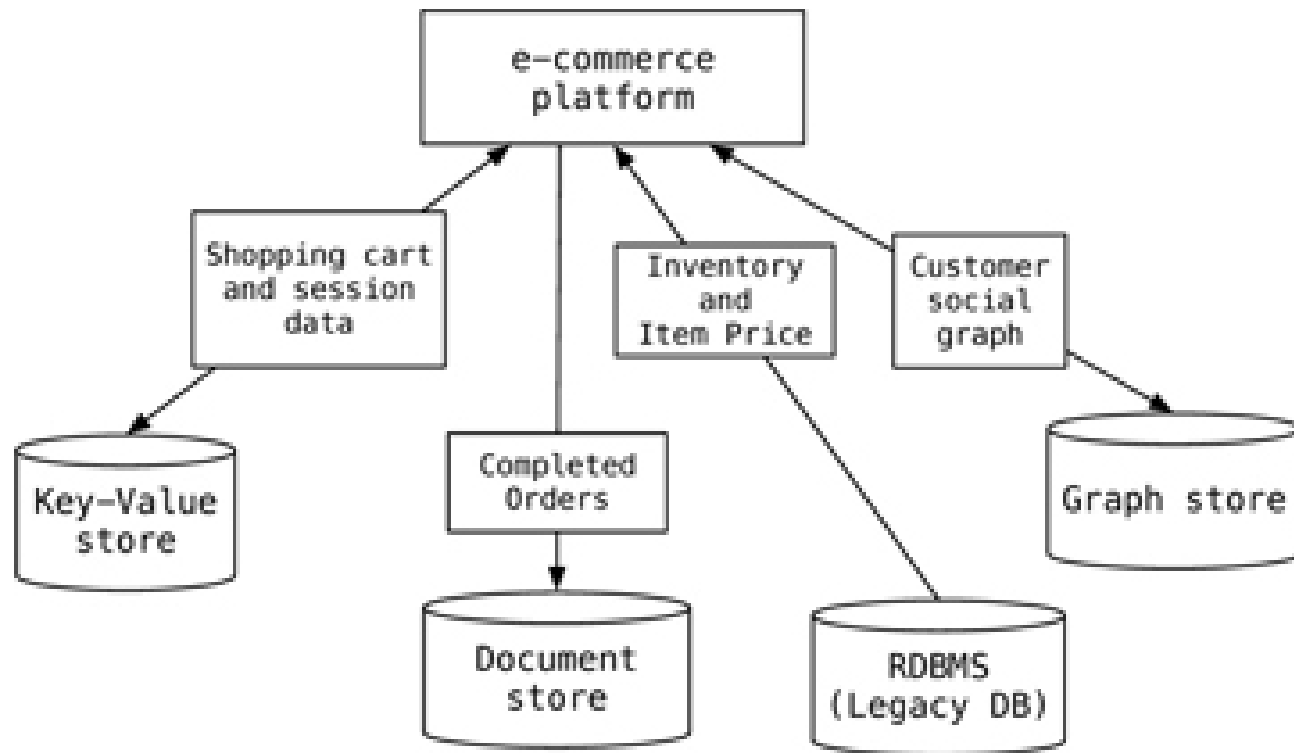



Figure 13.3. Example implementation of polyglot persistence



Martin Fowler

Relational is fantastic but ***Polyglot Persistence*** tells us to use the right data store for the challenge we are facing!

If we have ***highly connected data***, a RDBMS might not be the best too for the job!

A Property Graph











Entities become ***NODES***

FK/PK relationships become ***EDGES***

Many-To-Many relationships become ***EDGES***

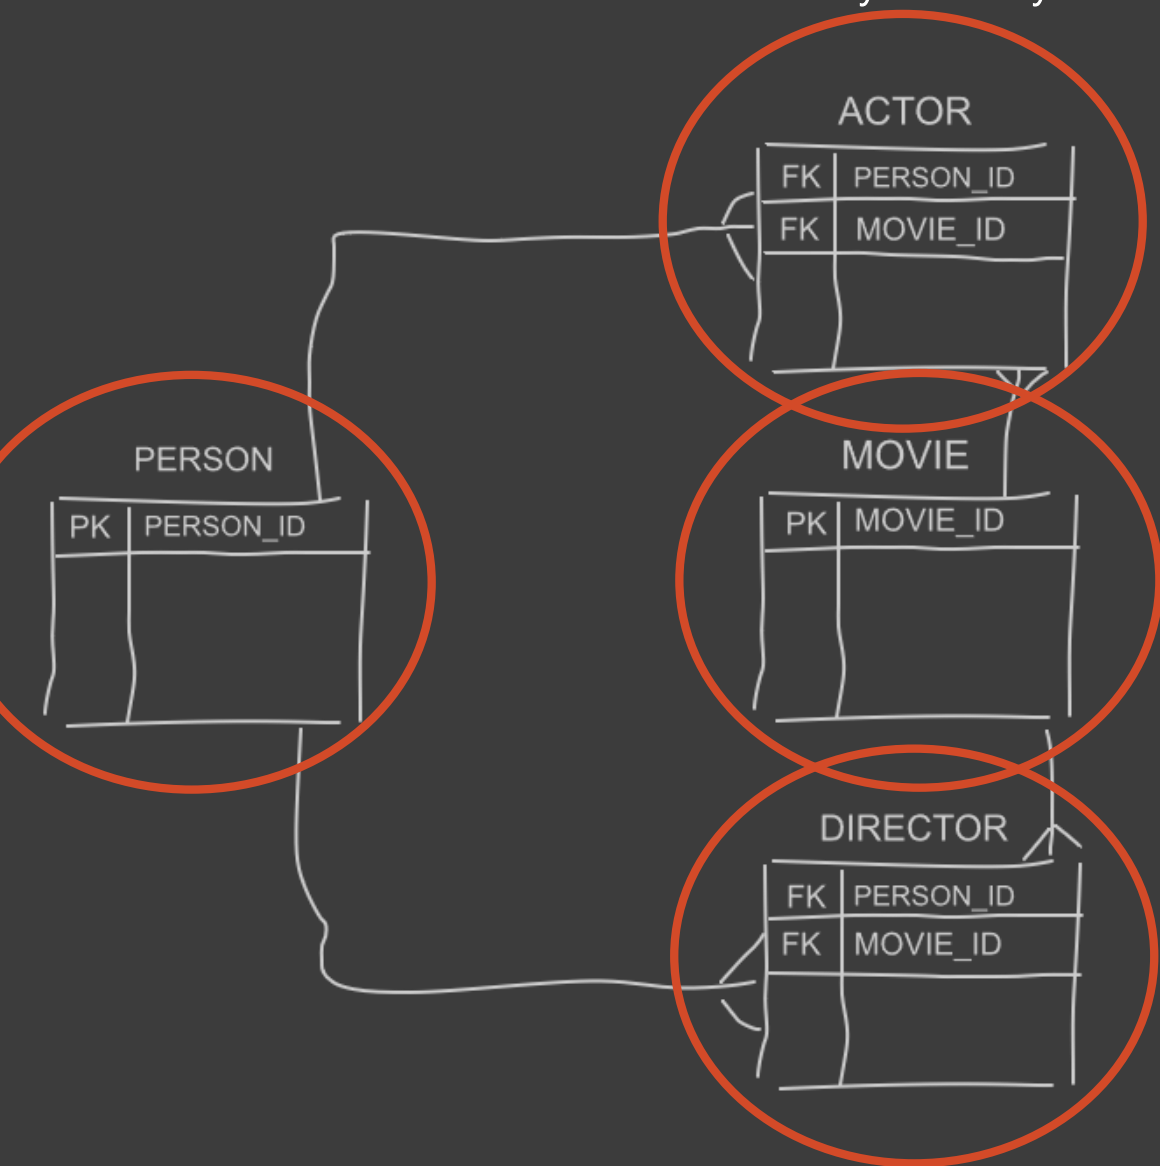
Attributed joins become ***EDGEs*** with properties

EDGE properties

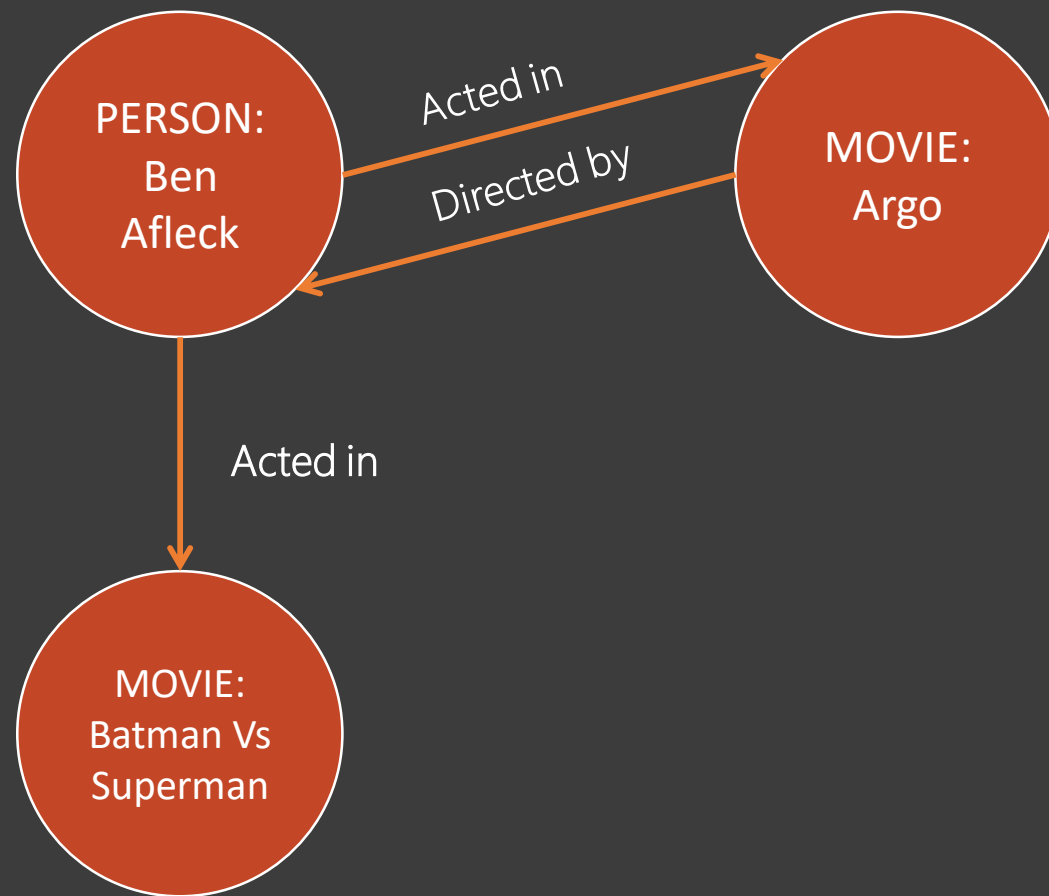




Many-To-Many



Many-To-Many



Entities become ***NODES***

FK/PK relationships become ***EDGES***

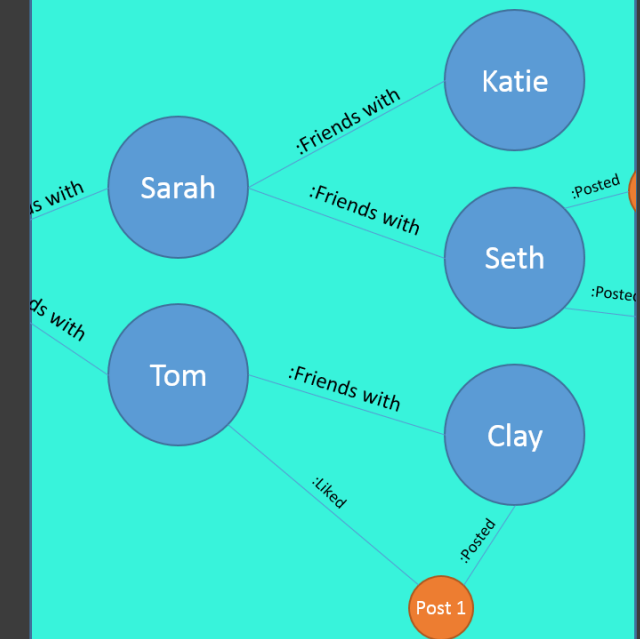
Many-To-Many relationships become ***EDGES***

Attributed joins become ***EDGEs*** with properties

1. **Shortest path** – What is the shortest path between 2 NODES
2. **Transitive closure** – Find a NODE n hops away – Find the NODE which is 10 hops away from where I am now
3. **Polymorphism** – Find any NODE connected to another NODE

Social Graph

Who knows who? Spot queen bees



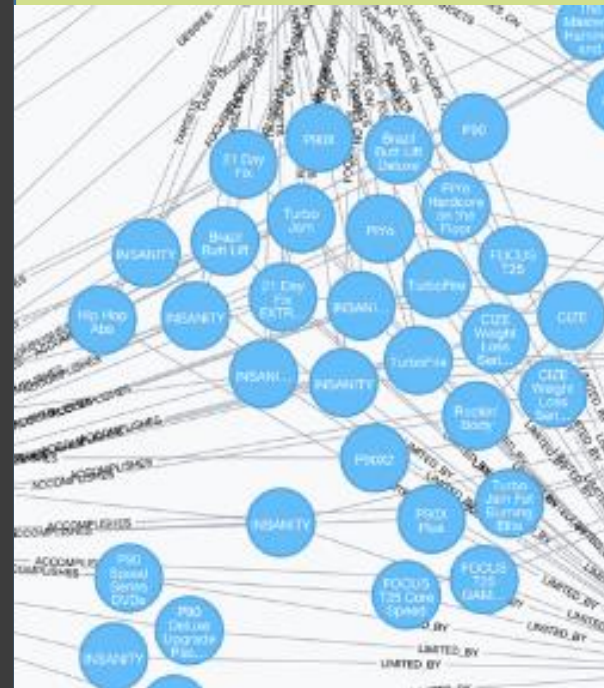
Fraud Detection

Spot rings of fraudulent activity



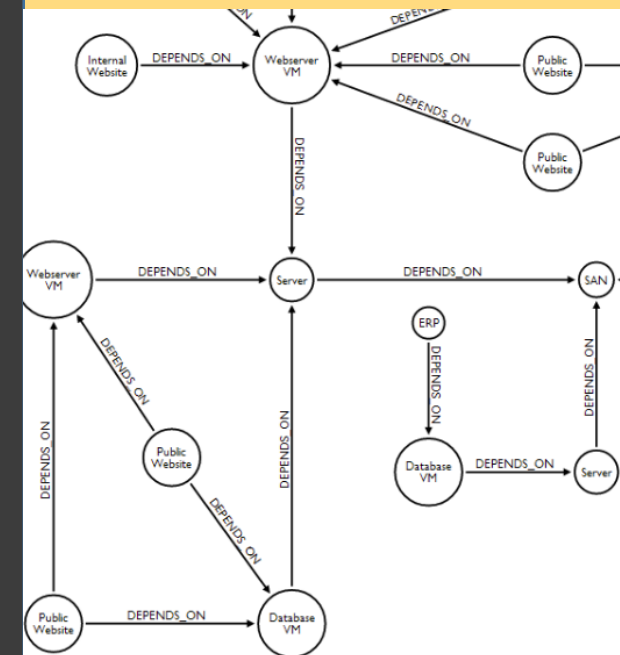
Recommendation

Bob bought X so did
Jim, maybe Jim
would also like...



Network & IT

Understand the
interconnectivity of
your network



How does this relate to SQL Server?



- Supports **Graph objects** & **Graph queries** to analyze complex relationships
- **Adaptive query processing** learns & optimizes for unparalleled performance
- **Python** and **R** support
- Advanced ML + **Deep Learning on GPUs**

Industry-leading performance on the most secure data platform, with built-in intelligence for all your data



NODE aka vertexes

```
DROP TABLE IF EXISTS Genre;  
CREATE TABLE Genre (GenreId INTEGER PRIMARY KEY, Genre VARCHAR(100)) AS NODE;
```

A node table represents an entity in a graph schema. Every time a node table is created, along with the user defined columns, an implicit \$node_id column is created, which uniquely identifies a given node in the database.

\$node_id - It is recommended that users create a **unique constraint or index** on the \$node_id column at the time of creation of node table, but if one is not created, a default unique, non-clustered index is automatically created.

EDGE aka relationship

```
DROP TABLE IF EXISTS ActedIn;  
CREATE TABLE ActedIn AS EDGE;
```

An edge table represents a **relationship** in a graph. Edges are always directed and connect two nodes. An edge table enables users to model many-to-many relationships in the graph.

Each edge table has three pseudo-columns:

\$edge_id: The id of the edge record

\$from_id: One of the nodes in the edge record

\$to_id: The other node in the edge record

MATCH

```
SELECT
    M.Movie
    , U1.FullName
    , U2.FullName
FROM dbo.[USER] U1, dbo.[USER] U2, dbo.Movie M, dbo.Reviewed R1,  dbo.Reviewed R2
WHERE MATCH (U1-(R1)->M<-(R2)-U2)
```

Specifies a search condition for a graph. MATCH can be used only with graph node and edge tables, in the SELECT statement as part of WHERE clause.

ASCII ART Syntax based on Cypher
Node-(edge)->Node OR Node<-(edge)-Node

Demo

Featured Dataset

TMDB 5000 Movie Dataset

Metadata on ~5,000 movies from TMDb



The Movie Database (TMDb) · last updated a month ago

[Overview](#)[Data](#)[Kernels](#)[Discussion](#)[Activity](#)[Download \(9 MB\)](#)[New Kernel](#)

Tags

film

medium

featured

Top Contributors

[Anisotropic](#)

1st

[FabienDaniel](#)

2nd

[AdhokshajaPradeep](#)

3rd

Kernels

[Principal Component Analysi...](#)

run a month ago

184

votes

[Film recommendation engine](#)

run 6 days ago

117

votes

[EDA with Plotly](#)

run a year ago

43

votes

Discussion

[Source for opening weeken...](#)

2 days ago

0

replies

[Updating datasets and deleti...](#)

10 days ago

1

reply

[Principal Component Analy...](#)

10 days ago

52

replies

Description

Background

What can we say about the success of a movie before it is released? Are there certain companies (Pixar?) that have found a consistent formula? Given that major films costing over \$100 million to produce can still flop, this question is more important than ever to the industry. Film aficionados might have different interests. Can we predict which films will be highly rated, whether or not they are a commercial success?

This is a great place to start digging in to those questions, with data on the plot, cast, crew, budget, and revenues of several thousand

Demo

Depth	RDBMS execution time(s)	Neo4j execution time(s)	Records returned
-------	-------------------------	-------------------------	------------------

2	0.016	0.01	~2500
---	-------	------	-------

MATCH caveats

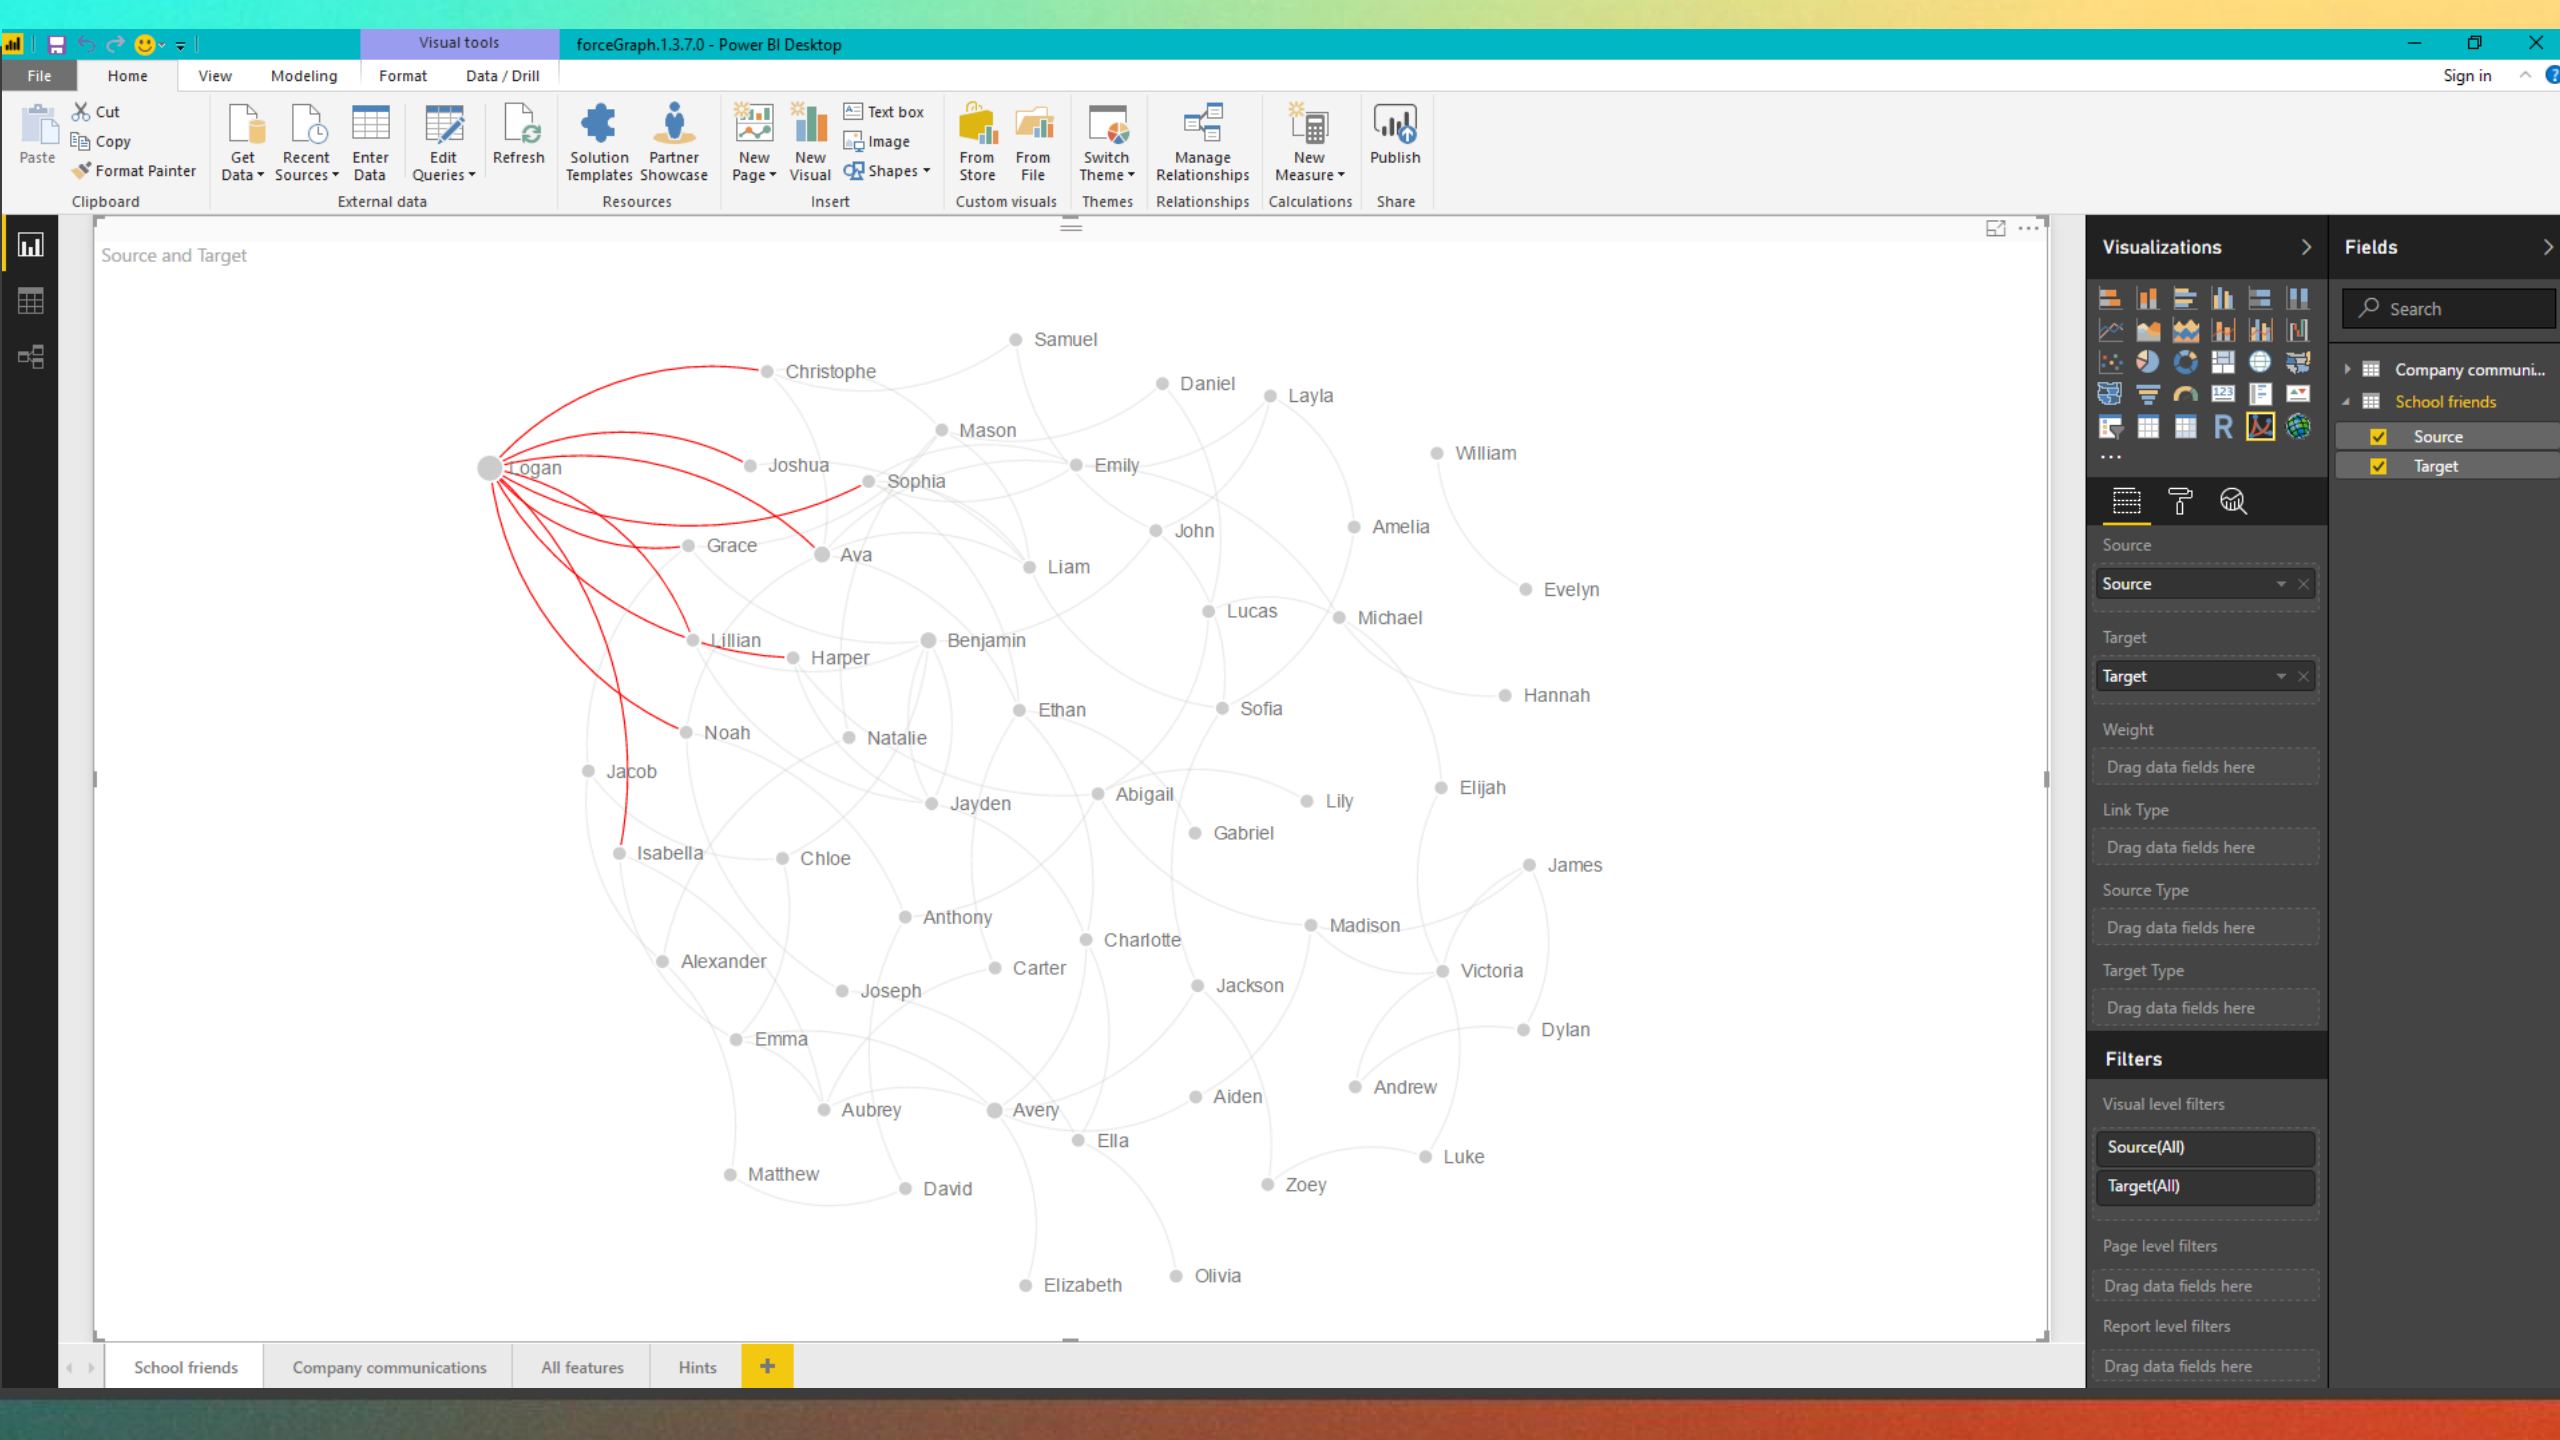
- The node names inside **MATCH can be repeated**. In other words, a node can be traversed an arbitrary number of times in the same query.
- An edge name **cannot be repeated** inside MATCH.
- An edge can point in either direction, but it must have an **explicit direction**.
- **OR and NOT** operators are **not supported** in the MATCH pattern. MATCH can be combined with other expressions using AND in the WHERE clause. However, combining it with other expressions using OR or NOT is not supported

Limitations

1. **Shortest path** – What is the shortest path between 2 NODES
2. **Transitive closure** – Find a NODE n hops away – Find the NODE which is 10 hops away from where I am now
3. **Polymorphism** – Find any NODE connected to another NODE

Limitations continued

- Local or global **temporary tables** cannot be node or edge tables.
- Table types and **table variables** cannot be declared as a node or edge table
- Node and edge tables cannot be created as **system-versioned temporal tables**.
- Node and edge tables cannot be **memory optimized tables**.
- Users cannot update the \$from_id and \$to_id columns of an edge using **UPDATE** statement. To update the nodes that an edge connects, users will have to insert the new edge pointing to new nodes and delete the previous one.
- **Cross database queries** on graph objects are not supported.



<Appendix> - Research

- <https://www.mssqltips.com/sqlservertip/4883/sql-server-2017-graph-database-example/>
- http://sqlblog.com/blogs/john_paul_cook/archive/2017/06/19/modeling-many-to-many-relationships-in-sql-server-2017-graph-database.aspx
- <https://stephanefrechette.com/sql-graph-sql-server-2017/#.WVbCP2grKUI>
- <https://blogs.technet.microsoft.com/dataplatforminsider/2017/04/20/graph-data-processing-with-sql-server-2017/>
- <https://stackoverflow.com/questions/20979831/recursive-query-used-for-transitive-closure>
- <https://blogs.msdn.microsoft.com/sqlcat/2017/04/21/build-a-recommendation-system-with-the-support-for-graph-data-in-sql-server-2017-and-azure-sql-db/>
- <https://github.com/arvindshmicrosoft/MillionSongDatasetinSQLServer>
- <https://github.com/Microsoft/sql-server-samples/blob/master/samples/applications/iot-connected-car/README.md>
- <https://www.youtube.com/watch?v=3vleFXDGoEs>
- <https://www.simple-talk.com/sql/t-sql-programming/sql-graph-objects-sql-server-2017-good-bad/>
- <https://sonra.io/2017/06/12/benefits-graph-databases-data-warehousing/>
- <https://social.technet.microsoft.com/wiki/contents/articles/37993.t-sql-graph-based-tables-in-sql-2017.aspx>

Graph processing is amazing!

However, SQL Server is not quite there yet! When Graph Processing in SQL Server can rival Neo4j, we will have a very powerful system.

Thank you

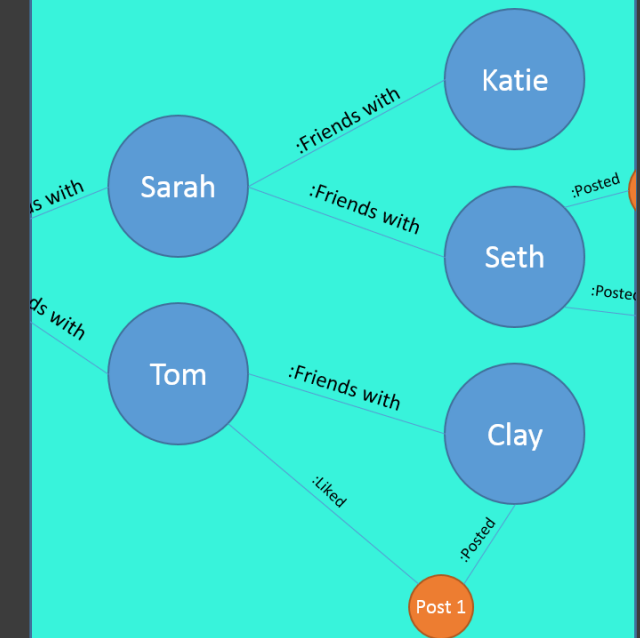


@SQLShark
tpm@adatis.co.uk
www.adatis.co.uk
@AdatisBI

So we have extra time...

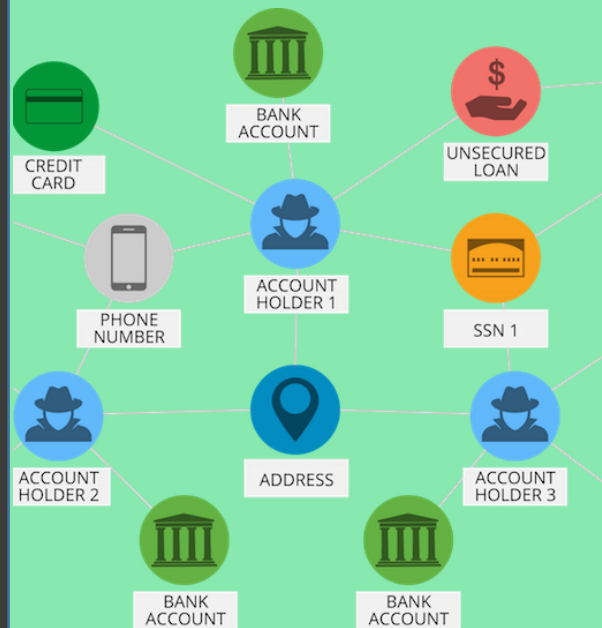
Social Graph

Who knows who? Spot queen bees



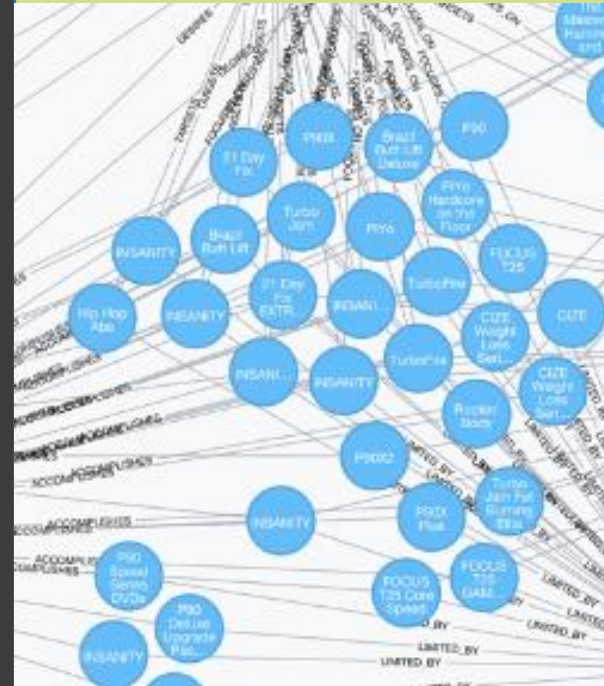
Fraud Detection

Spot rings of fraudulent activity



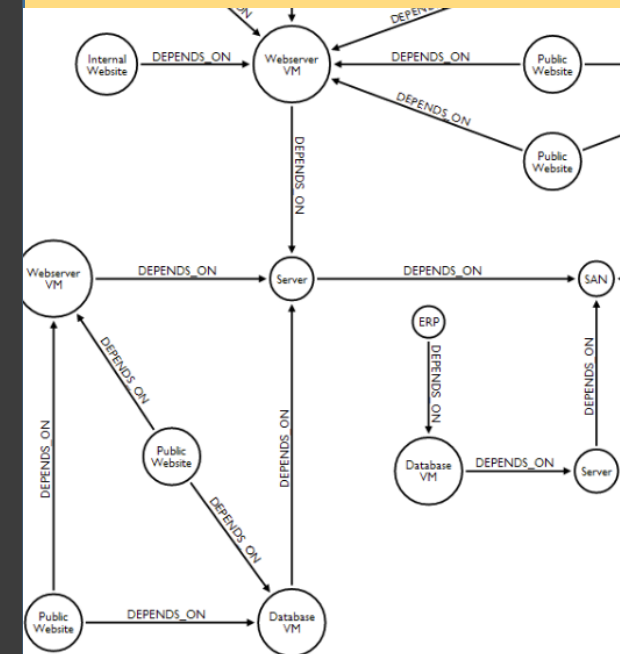
Recommendation


Bob bought X so did
Jim, maybe Jim
would also like...

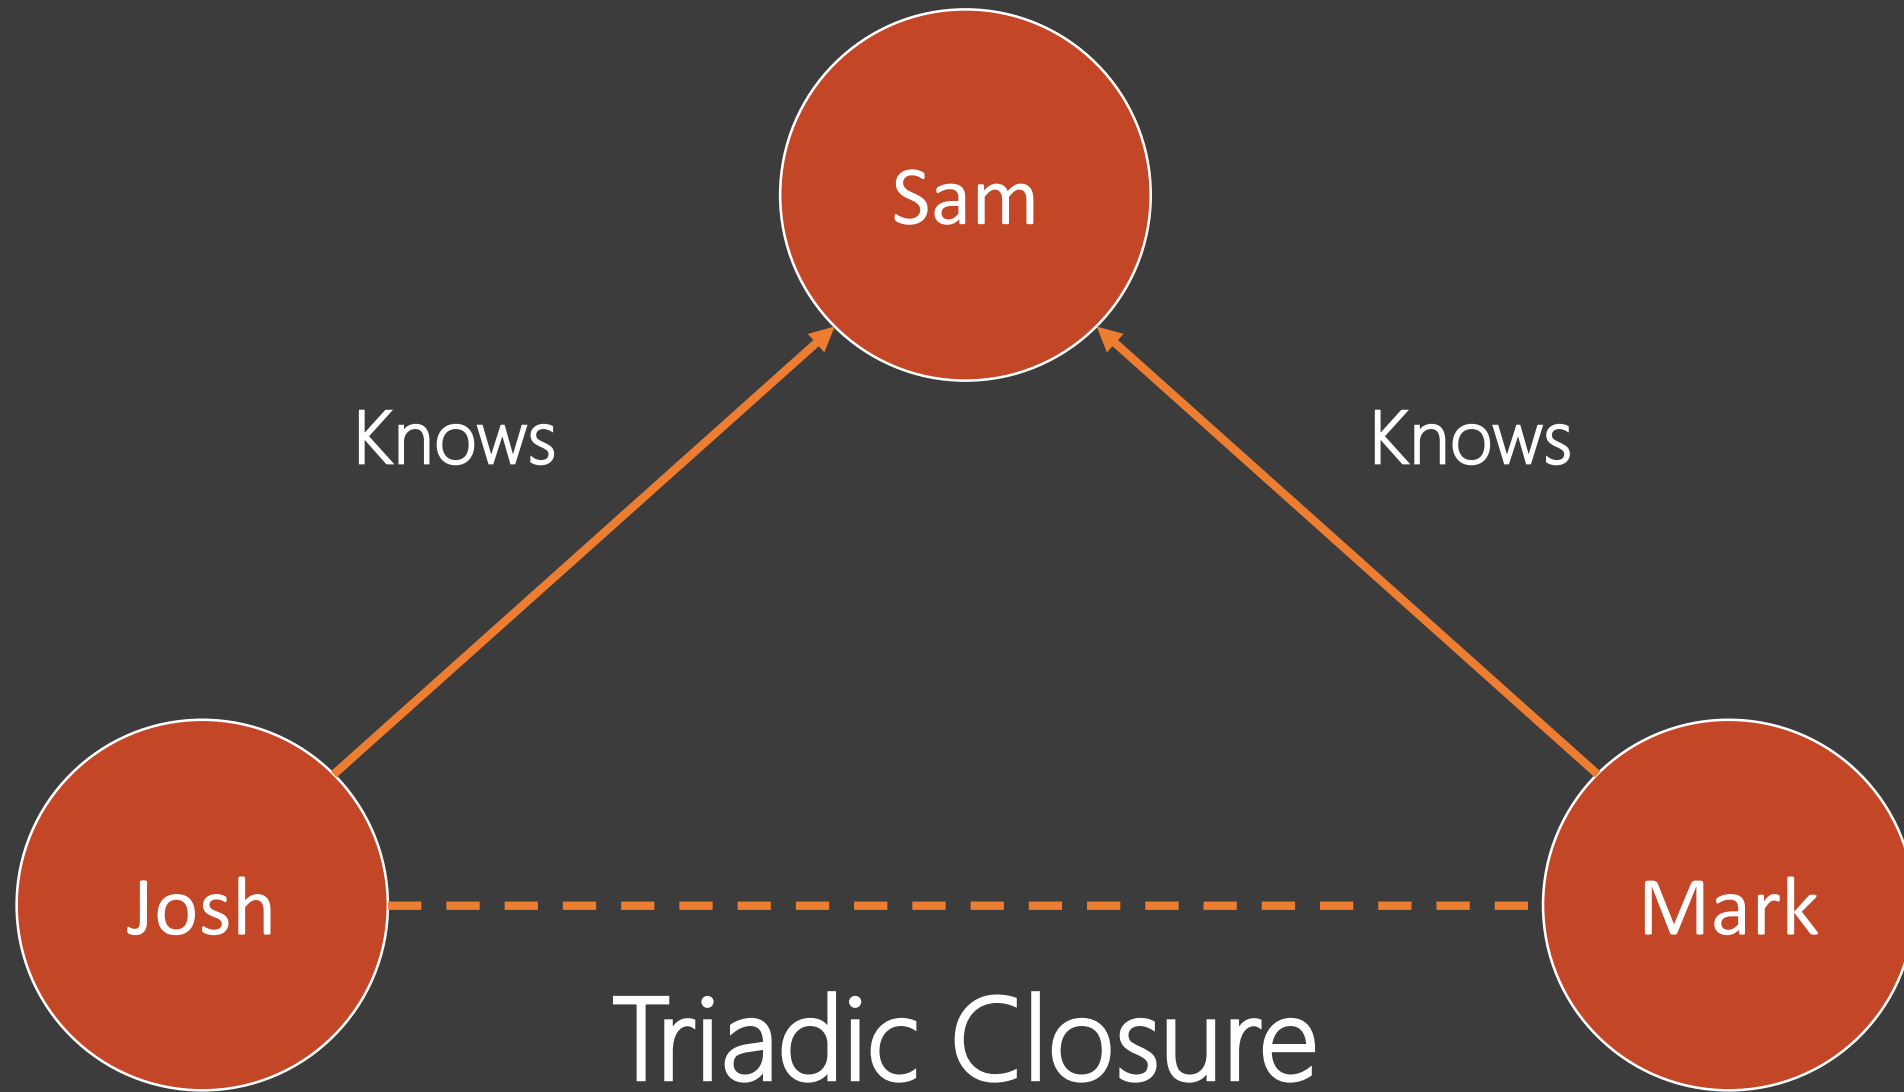


Network & IT

Understand the
interconnectivity of
your network







Triadic Closure

Recommendation Demo

